


# Weather Research & Forecasting Model (WRF)

## Processing of forecasts with docker containers in Azure cloud environment



Ricardo Faria  
Luís Fernando Nóbrega



# Environment setup

## Important aspects

- ❖ Used distribution: **ubuntu** 22.04;
- ❖ WRF Install Script:  
<https://github.com/bakamotokatas/WRF-Install-Script/blob/master/README.md>;
- ❖ Install script: WRF4.6.0\_Install.bash;
- ❖ In order to save space while maintaining resolution, **GEOGRID.TBL** was altered to use 30s res as default and worse resolutions such as 2m, 5m, 10m , 1 deg and 2 deg were manually **deleted**;
- ❖ Image name: **reduced\_ubuntu\_image**;
- ❖ Image disk usage: 32.4 GB;
- ❖ tar.gz disk usage: 3.3 GB;
- ❖ Uncompressing time: Around **15 min** for decompression of \*tar.gz

# Docker container

## Docker file structure:

**FROM** ubuntu:22.04

**LABEL** Ricardo Faria <email@gmail.com>

**ENV** DEBIAN\_FRONTEND=noninteractive \  
LANG=en\_US.UTF-8 \  
LANGUAGE=en\_US:en \  
LC\_ALL=en\_US.UTF-8

**RUN** apt-get update && \  
apt-get install -y \  
sudo \  
file \  
nano \  
locales \  
python3 \  
hostname \  
m4 \  
make \  
perl \  
tar \  
bash \  
tcsh \  
time \  
wget \  
cmake \  
pkg-config \  
libxml2-dev \  
libcurl4-openssl-dev \  
libnetcdf-dev && \  
apt-get clean && \  
rm -rf /var/lib/apt/lists/\* && \  
locale-gen en\_US.UTF-8 && \  
update-locale LANG=en\_US.UTF-8 LC\_ALL=en\_US.UTF-8 && \  
useradd -m swe && echo "swe ALL=(ALL) NOPASSWD:ALL" \  
> /etc/sudoers.d/swe && \  
echo "alias cp='cp -iv'" >> /home/swe/.bashrc && \  
echo "alias mv='mv -iv'" >> /home/swe/.bashrc && \  
echo "alias mkdir='mkdir -pv'" >> /home/swe/.bashrc && \  
echo "alias ll='ls -FGIAhp'" >> /home/swe/.bashrc

**USER** swe  
**WORKDIR** /home/swe  
**CMD** ["bash"]

Only one **RUN** to  
reduce number of  
layers and  
occupied space

## Other actions performed after docker initiation

- ❖ apt install pip
- ❖ pip install requests
- ❖ Manually set In -s ungrib/Variable\_Tables/Vtable.GFS Vtable for first time
- ❖ Same for ls -ls geogrid/GEOGRID.TBL

## Resolution limitations (30s)

- ❖ **Deleted:** varssso\_10m, varssso\_5m, varssso\_2m, orogwd\_2deg, orogwd\_1deg, orogwd\_30m, orogwd\_20m, lai\_modis\_10m;
- ❖ **Kept:** albedo\_modis      lai\_modis\_30s modis\_landuse\_20class\_30s\_with\_lakes  
soiltemp\_1deg      soiltype\_top\_30s      varssso greenfrac\_fpar\_modis maxsnowalb\_modis  
orogwd\_10m      soiltype\_bot\_30s topo\_gmted2010\_30s

To add other resolutions: [source1](#) or [source2](#) in WPS\_GEO directory

# Docker container commands

<code>docker images</code>	Lists all available images, size, ID, creation date and repository
<code>docker ps -a</code>	Lists all containers, exit status, creation date and container ID
<code>docker run -it my_container_name</code>	Starts a container in interactive mode. Only practical way to alter container files and internal architecture
<code>docker commit container_id image_name</code>	Allows altering of base image after exiting from container and retrieving container ID. Without it, changes remain local inside container
<code>exit</code> (from inside container)	To exit container from interactive mode. Should be followed by <code>docker ps -a</code> to check status of exit, which should be 0 for saves to be made
<code>docker system df</code>	Displays disk usage from whole docker environment
<code>docker image prune</code>	Removes dangling image creation layers -> <b>USE WITH CAUTION</b>
<code>docker stop my_container</code>	Stops containers running (from the outside)
<code>docker remove my_container</code>	Removes stopped containers

# Editing the container interactively

```
fernando@DESKTOP-L4H7EVR:~/OOM/copernicus3_images$ docker images 1 - choose image
REPOSITORY          TAG         IMAGE ID      CREATED       SIZE
reduced_ubuntu_image latest      0c67e48bfa8b 6 days ago   32.4GB
my-ubuntu-22.04-image latest      db73bb86acd9 10 days ago  65.9GB

fernando@DESKTOP-L4H7EVR:~/OOM/copernicus3_images$ docker run -it reduced_ubuntu_image 2 - run interactive mode
swe@126a61b4df09:~$ cd Build_WRF/
swe@126a61b4df09:~/Build_WRF$ ls
CONFIGS  LIBRARIES  WPS-4.6.0  WPS_GEOG  WRF-4.6.0-ARW  3 - navigate dirs and change what is needed
swe@126a61b4df09:~/Build_WRF$ cd CONFIGS/
swe@126a61b4df09:~/Build_WRF/CONFIGS$ ls
forecast_download.py  forecast.sh  historic_download.py  instructions.txt  model_set.py  namelist_editor.py  processors.py  wps_input.txt  wrf_input.txt
swe@126a61b4df09:~/Build_WRF/CONFIGS$ exit 4 - leave interactive mode
exit
fernando@DESKTOP-L4H7EVR:~/OOM/copernicus3_images$ docker ps -a 5 - see available containers to fetch ID
CONTAINER ID   IMAGE          COMMAND                  CREATED        STATUS              PORTS          NAMES
126a61b4df09   reduced_ubuntu_image "bash"                 18 seconds ago Exited (0) 3 seconds ago             inspiring_chaplygin

fernando@DESKTOP-L4H7EVR:~/OOM/copernicus3_images$ docker commit 126a61b4df09 my-ubuntu-22.04-image 6 - fetch changes to new image
```

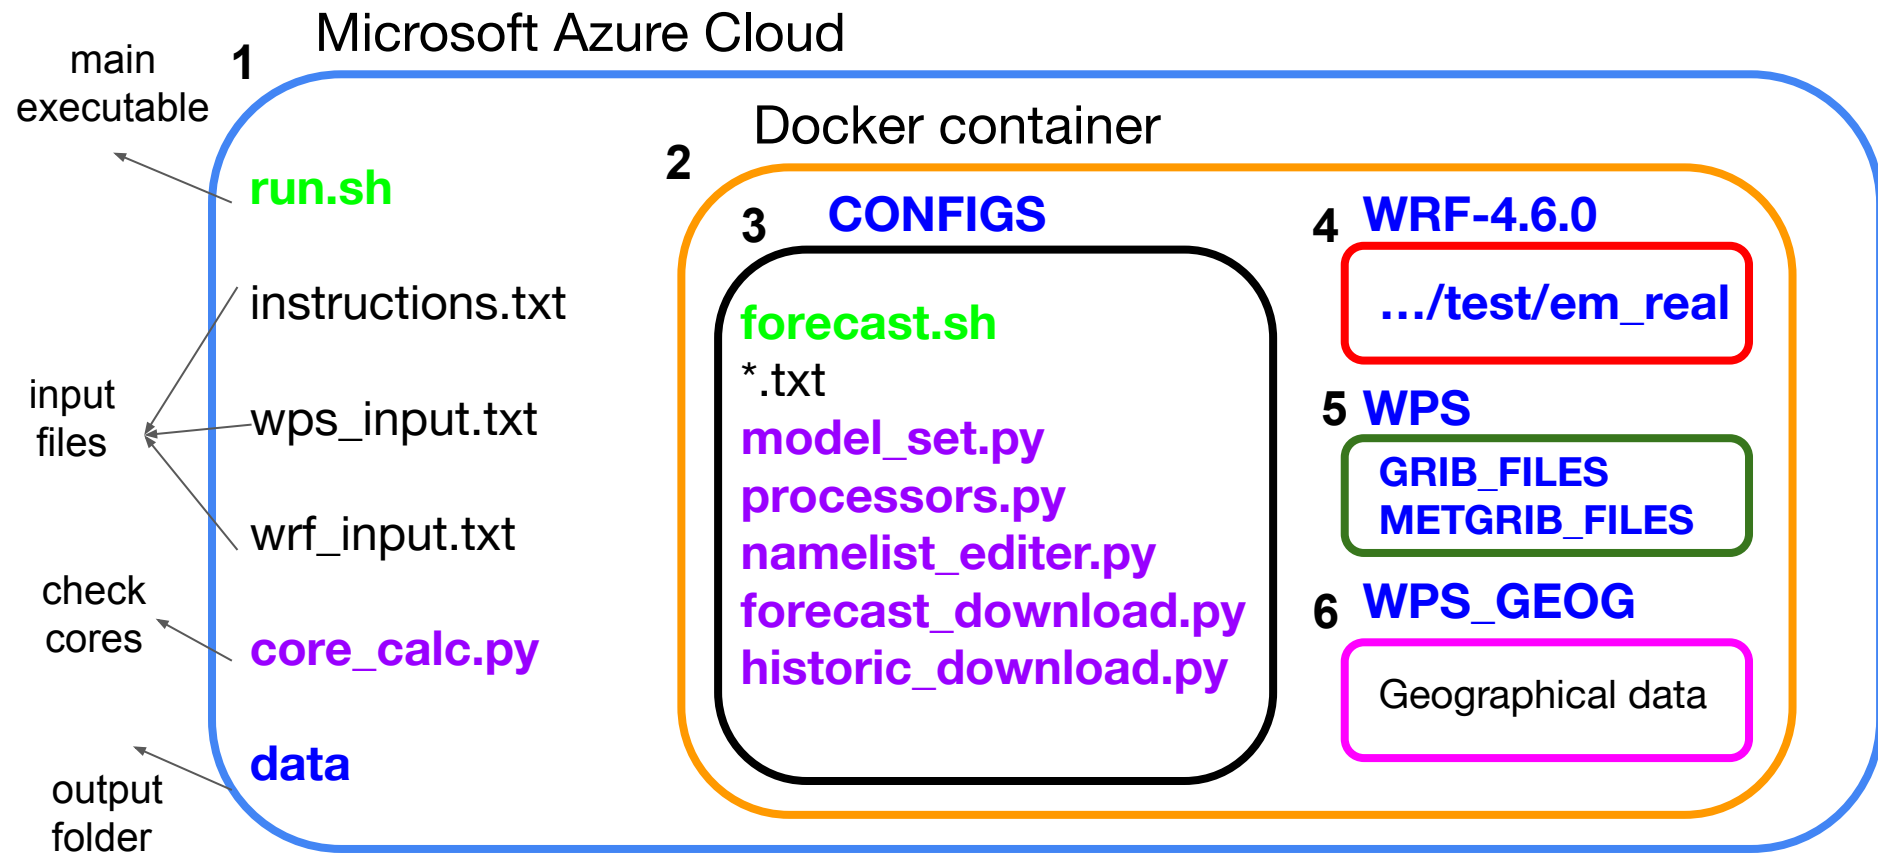
An id: usually something like sha:12@asd4\$%39 will appear -> use `$ docker images` to check for updated date

If docker seems to take too much space:

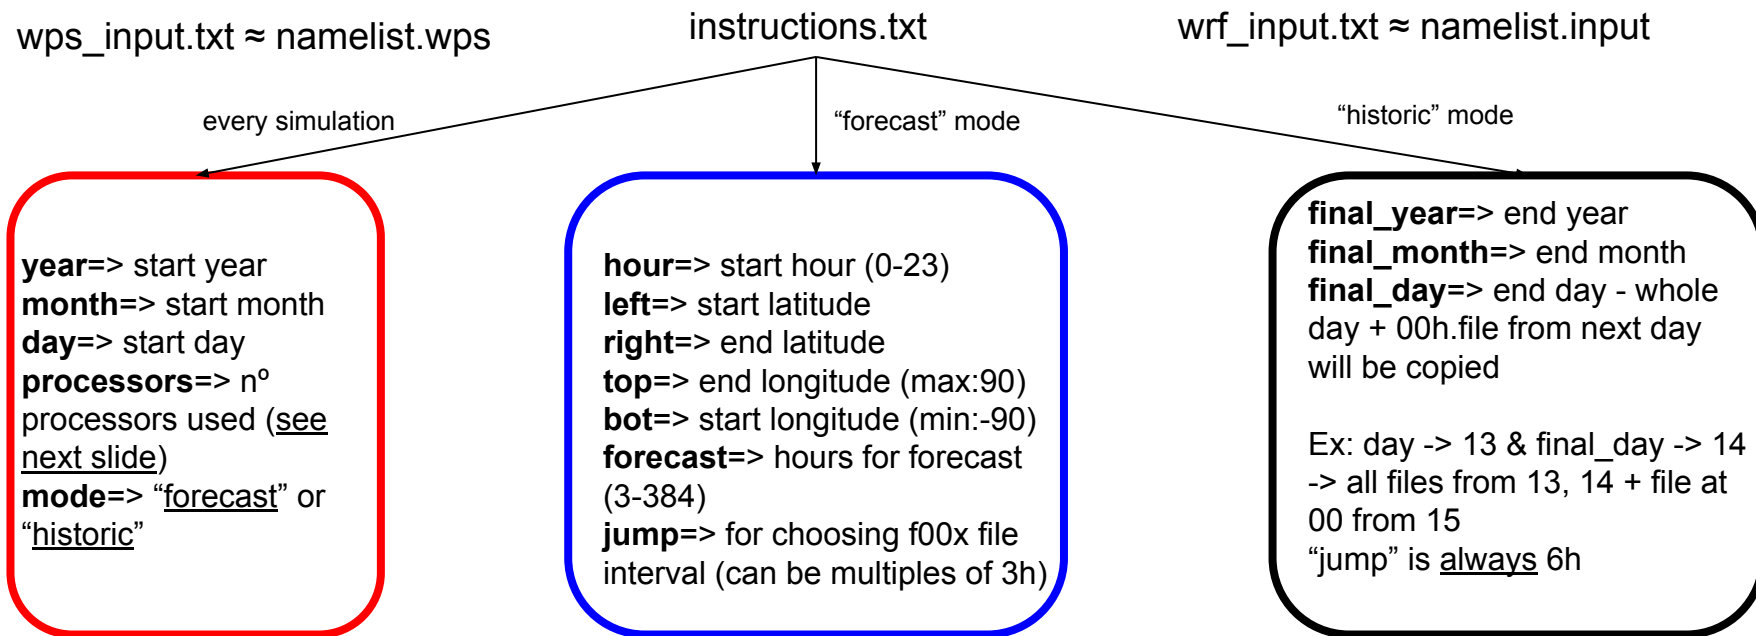
**DON'T** use `-a` flag after `$ docker image prune`  
it deletes images without associated  
containers

```
fernando@DESKTOP-L4H7EVR:~/OOM/copernicus3_images$ docker system df
TYPE          TOTAL          ACTIVE        SIZE          RECLAIMABLE
Images        2              0             98.29GB       98.29GB (100%)
Containers    0              0             0B            0B
Local Volumes 0              0             0B            0B
Build Cache   23            0             446.7MB       446.7MB
fernando@DESKTOP-L4H7EVR:~/OOM/copernicus3_images$ docker image prune
```

# Running Environment structure



## Necessary input files before **./run.sh**



**Notes:** The simulation working directory **MUST** have a instructions.txt similar to the provided one .The files should be manually revised before **./run.sh** .All terms are **CASE SENSITIVE** and should not have any capitalization. **DO NOT** change the **Bold** terms.



## Example input.txt

```
##### READ ME
# forecast data from "Data Transfer: NCEP GFS Forecasts (0.25 degree grid)" -> Data until 10days back
# most recent forecast data available -> 6h before current date (updated at 6,12,18,24 h) -> max forecast -> 384h -> aprox 14days
#
# historic data from "NCEP GDAS/FNL 0.25 Degree Global Tropospheric Analyses and Forecast Grids" -> Data until 10 years back
# most recent historic data available -> 48h before current date
##### For both setups (historic or forecast)
```

```
year=2020
month=8
day=13
```

```
processors=4
mode=historic
```

```
##### For forecast
# left & right -> latitude (min:0 max:360)
# top & bot -> longitude (min:-90 max:90) -> top > bot
```

```
hour=0
left=250
right=360
top=40
bot=-5
forecast=3
jump=3
```

```
##### For historic
# final_day -> whole day + 00h.file from next day will be copied
# Ex: day -> 13 & final_day -> 14 -> all files from 13, 14 + file at 00 from 15
# for data of just 1 entire day -> day same as final_day
# hour_jump is always 6h in this mode
```

```
final_year=2020
final_month=8
final_day=13
```

For multiple domains, the simulation period will be the same. Coordinates should accommodate the largest domain

These must include **ref\_lat** and **ref\_lon** that is in wps\_input.txt

This file changes automatically, however **make sure this exists** before running anything

## Choosing appropriate number of processors

**core\_calc.py** -> interactive function to check appropriate number of cores to be used. Checks wps\_input.txt and retrieves a core interval. It **DOES NOT** change the value in instructions.txt but stops program in case of excessive nº cores -> returns MIN processors: # and MAX processors: # as well as a **warning** if size and time parameters of **first domain** do not match in wrf\_input.txt, wps\_input.txt and instructions.txt

For your **smallest-sized** domain:

$((e\_we)/25) * ((e\_sn)/25)$  = most amount of processors you should use

For your **largest-sized** domain:

$((e\_we)/100) * ((e\_sn)/100)$  = least amount of processors you should use

If simulation is **failing**, usually use **more** processors!

For more information: [Number of cores](#)

**core\_calc.py**  
always runs  
before **run.sh**  
to prevent  
mismatches or  
missing files

## Starting Azure VM instance

```
# Log in to Azure
az login

# Create a resource group if you don't have one
az group create --name myResourceGroup --location eastus

# Create the virtual machine
az vm create \\  
  --resource-group myResourceGroup \\  
  --name myVM \\  
  --image UbuntuLTS \\  
  --size Standard_D32ls_v5 \\  
  --admin-username azureuser \\  
  --generate-ssh-keys
```

```
# Save the Docker image locally
docker save -o mydockerimage.tar mydockerimage:latest

# Transfer the image to the Azure VM
scp mydockerimage.tar azureuser@<vm-public-ip>:/home/azureuser/
```

↓

```
docker run -d --name mycontainer -e START_DATE=2023-01-01 -e END_DATE=2023-12-31 \\  
  -v /mnt/data:/app/data myContainerRegistry.azurecr.io/mydockerimage:latest
```

↓

```
# Stop the VM
az vm stop --resource-group myResourceGroup --name myVM

# Optionally, delete the VM and associated resources
az vm delete --resource-group myResourceGroup --name myVM --yes --no-wait
```

Check the Logs:

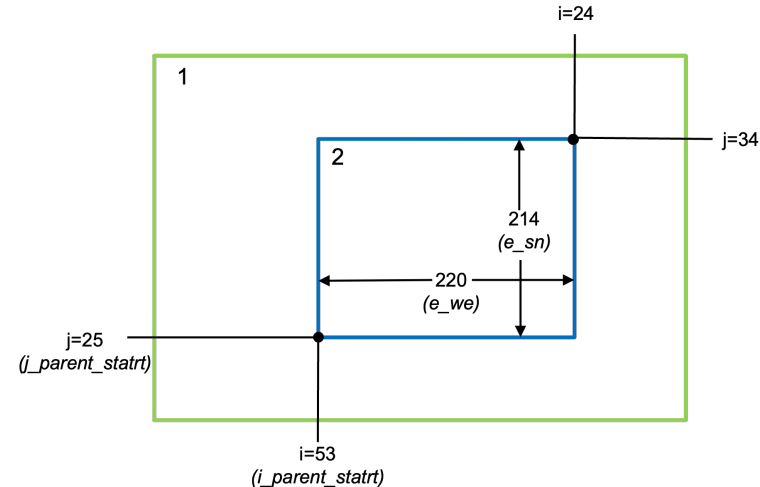
```
docker logs mycontainer
```

Download Data:

```
scp azureuser@<vm-public-ip>:/path/to/data /local/path/
```

## Notes about Nested Domains

- ❖ Geogrid is very sensitive to entry params;
- ❖ N° domains is set with **max\_dom** = 1;
- ❖  $e_{we} = 100,46$ , and  $e_{sn} = 100,37$ ,
- ❖  $(e_{sn}-s_{sn}+1) = m \times \text{parent\_grid\_ratio} + 1$  (m is an integer);
- ❖ In wrf.input **topo\_wind** = 1, 0, 0, might need to be set to 0 in nested domains due to insufficient data;
- ❖ **i\_parent\_start** and **j\_parent\_start** define the domains position. If nested domain is outside parent domain, make sure to reduce size or change the i and j start tiles.
- ❖ Alternatively, increase **parent\_grid\_ratio**



## Running command

```
./run.sh -e START_DATE=2023-01-01 -e END_DATE=2023-01-02
```

- ❖ START\_DATE and END\_DATE **must** be provided even if files are correctly set;
- ❖ Both environment variables can be in %Y-%m-%d or %Y-%m-%d\_%H:%M:%S;
- ❖ Even though the format asks for minutes and seconds, they have to be 00:00 as the data source only contains hourly files.
- ❖ By default, if %Y-%m-%d is provided, it will auto fill it with 00:00:00;
- ❖ Before proceeding a input will be required -> **CHECK** everything before running.
- ❖ **Warnings** might be given -> Some are normal depending on the format of instructions;
- ❖ Ex: WARNING: wps and wrf input files do not have the same dx -> if multiple domains with different sizes are used, this will appear for sure;
- ❖ For historic mode, choose periods that are multiples of 24h.

# Docker container

**forecast.sh ->** has all instructions for executing the remaining files, geogrid, ungrib, metgrib, run.exe and wrf.exe. Almost error messages related with missing files or subprocesses are here. Check example file for unbugging.

**\*.txt ->** all txt files that were copied from Azure cloud and will change inside files in WPS and WRF dirs

**model\_set.py ->** reads instructions.txt to determine which mode to run, “forecast” or “historic” -> case sensitive

**processors.py ->** reads instructions.txt to determine number of processors to use -> case sensitive

**namelist\_editor.py ->** changes namelist.input and namelist.wps with the .txt provided in Azure cloud

**forecast\_download.py ->** checks connection and downloads grib files for “forecast” mode. Will search based on terms such as “top”, “bot”... etc -> IF download fails, either URLS changed and file must be changed manually or too big of a domain was chosen. If grib files exceed 1GB, URL request error may occur -> check log for errors

[https://nomads.ncep.noaa.gov/cgi-bin/filter\\_gfs\\_0p25\\_1hr.pl](https://nomads.ncep.noaa.gov/cgi-bin/filter_gfs_0p25_1hr.pl)

**historic\_download.py ->** checks connection and downloads fnl files for “historic” mode. Unlike forecast, the files will always include the whole world and not a snippet domain. <https://rda.ucar.edu/datasets/d083003/dataaccess/#>

# Debugging and common errors

- ❖ **forecast.sh** instantly **failed** -> incorrect format for input files
- ❖ **Download failed or taking too long** -> domain choice too big; grib files may be too big; URL provider might have changed (check demo links); time period might be too big; dates in instructions.txt might not be correct; Check if wrong operating mode was chosen “forecast” or “historic”; Internet connection might be too slow.
- ❖ **Geogrid failed** -> incorrect info in wps\_input (maybe in instructions.txt)
- ❖ **Ungrib failed** -> incorrect info in wps\_input and instructions.txt; dates not coincident; date in instructions does not have all necessary points referenced in wps\_input -> missing files; files might be too small due to download failure -> check size with “ll” in bash; files might be incomplete for performed tasks; with multiple domains, remember that the start and end dates must be the same
- ❖ **Metgrib failed** -> incorrect info in wrf\_input and wps\_input -> check compatibility and recheck instructions.txt; historic files might not contain all needed info for metgrid if they are too old; Chosen domain and coordinates might not be the same as the wps\_input ones -> check if ref\_lat and ref\_lon are within domain chosen for download data
- ❖ **Run.exe or wrf.exe failed** -> Insufficient data; Invalid time period etc... -> eventual failures usually occur before this stage
- ❖ **files not saved to /data** -> check console log for wrf.exe errors; chosen simulation time period might be too small

For further information about the process: [Compilation](#); [running\\_WPS](#); [WPS\\_namelist\\_variables](#); [running\\_WRF](#); [general\\_info](#)